

Knowledge Extraction & Retrieval

Step	Task	Tools / Models	
1. Data Ingestion	Text Documents	<ul style="list-style-type: none"> • PyMuPDF • PDFPlumber • Apache Tika • Azure Form Recognizer • Google Document AI • GPT-4.1 / GPT-4o] • Google Gemini 2.0 Flash / Pro 	
	Image-Based Documents	<ul style="list-style-type: none"> • Tesseract • LayoutLM • TrOCR • EasyOCR • PaddleOCR • Donut • Google Vision OCR • Microsoft Azure OCR • Claude 3.5 Sonnet 	
	Internal System Text	Working on it	
2. Preprocessing	<ul style="list-style-type: none"> • Cleaning & Normalization 	<ul style="list-style-type: none"> • spaCy • NLTK 	
3. Domain Understanding	<ul style="list-style-type: none"> • Named Entity Recognition (NER) • Key Phrase & Topic Extraction • Relation Extraction • Knowledge Graph Construction 	<ul style="list-style-type: none"> • spaCy custom NER • HuggingFace transformers • Prodigy (annotation tool) 	<ul style="list-style-type: none"> • RoBERTa-NER • BioBERT (medical NER) • FinBERT (finance NER) • med7 (medical entities) • BERT-base multilingual NER

4. Embedding Generation	<ul style="list-style-type: none"> Generate Text Embeddings 		<ul style="list-style-type: none"> OpenAI text-embedding-3-large / 3-small Llama 3 Instruct Models e5-large-v2 all-mpnet-base-v2
5. Vector Database	<ul style="list-style-type: none"> Store Text Embeddings 	<ul style="list-style-type: none"> Pinecone Qdrant Milvus Redis Vector Store Weaviate ChromaDB 	
6. Retrieval System	<ul style="list-style-type: none"> Semantic Search Hybrid Search (Keyword + Semantic) Contextual Retrieval (RAG) 	<ul style="list-style-type: none"> LangChain retrievers LlamaIndex Qdrant/Pinecone query APIs Elasticsearch Hybrid Search 	<p>ReRankers:</p> <ul style="list-style-type: none"> Cohere ReRank bge-reranker-large cross-encoder/ms-marco-MiniLM
7. Content Generation (LLM Layer)	<ul style="list-style-type: none"> Document Summarization Q&A From Knowledge Base Drafting New Content Multi-step Reasoning 		<ul style="list-style-type: none"> GPT-4/5 GPT-4o Llama 3 70B Claude 3.5 Sonnet
8. Continuous Learning & Feedback Loop	<ul style="list-style-type: none"> Reinforcement learning from human feedback Data augmentations Active learning Flow improvements 		<ul style="list-style-type: none"> Periodic fine-tuning of BERT/LLM models Lightweight adapters (LoRA/QLoRA)